



Figure 1: Comparison between the mirror map learned on Hopper and the negative entropy on three MuJoCo environments. The plot present the performance of PMD for both mirror maps, reporting the average over 8 realizations and a shaded region denoting the standard error around the average.